

Reinforcement Learning

COMS4047A/COMS7053A

Lab 3 - Deep Q-Network

Tamlin Love (tamlin.love1@students.wits.ac.za)
Shahil Mawjee (shahil.mawjee@students.wits.ac.za)

Overview

The lab will focus on the implementation of the Deep Q-Learning (DQN) agent to play Atari games. For this lab we will use the Pong ([link](#)) environment which are contained within the OpenAI Gym **Atari** package.

Goals:

- Implement algorithm presented in research paper.
- Understand hyper-parameter optimisation for Deep RL.

Compute

Training DQN agents can be computationally intensive, so we suggest using Google Colab if you do not have access to an Nvidia GPU. Ensure you attach the notebook to a kernel with GPU access, this can be done by the following:

- Select *Runtime* in toolbar
- then select *Change runtime type*
- and then select **GPU** as *Hardware accelerator*

Refer to <https://colab.research.google.com/notebooks/intro.ipynb> for more information about Google Colab.

Submission

Work in groups of 3 to 4 people. List group members in the *group_members.txt* file.

For this lab, we have provided starter-code which makes use of Pytorch, so you will need to edit it if you plan to use TensorFlow.

You do not have to use the starter code and can implement your DQN from scratch. If you are using Google Colab, save your Colab notebook as a notebook (.ipynb) file and include in your submission.

Include the following in your submission (as zip file):

1. A summary of the DQN algorithm. Make sure to include the following points:
 - Explain the training process outlined in figure 3.
 - Q-Network architecture.
 - List of hyper-parameters, provide the following for each parameter:
 - Name of parameter
 - Value used
 - Brief description of the role of the parameter
2. Your code files and saved model.
3. A plot of training reward curves, as shown in the paper.
4. A plot the loss curve during training.
5. A GIF of your trained DQN agent playing Pong.

Submit your zip file to Moodle.

1 Deep Q-Learning (DQN)

Read to the following papers in order to understand and implement the DQN:

- Human-level control through deep reinforcement learning (link)
- Playing Atari with Deep Reinforcement Learning (link)




Figure 1: Deep Q Network (Source: Nature paper)

Algorithm 1 Deep Q-learning with Experience Replay

```

Initialize replay memory  $\mathcal{D}$  to capacity  $N$ 
Initialize action-value function  $Q$  with random weights
for episode = 1,  $M$  do
  Initialise sequence  $s_1 = \{x_1\}$  and preprocessed sequenced  $\phi_1 = \phi(s_1)$ 
  for  $t = 1, T$  do
    With probability  $\epsilon$  select a random action  $a_t$ 
    otherwise select  $a_t = \max_a Q^*(\phi(s_t), a; \theta)$ 
    Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$ 
    Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$ 
    Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $\mathcal{D}$ 
    Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $\mathcal{D}$ 
    Set  $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$ 
    Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$ 
  end for
end for

```

1.1 Environment: Atari Pong

For this lab we will use Atari's Pong game ([link](#)).

```
pip install gym[atari]
```

In this environment, the observation is an RGB image of the screen, which is an array of shape (210, 160, 3)




Figure 2: Pong Game (Source - OpenAI Gym website)

1.2 Training Loop

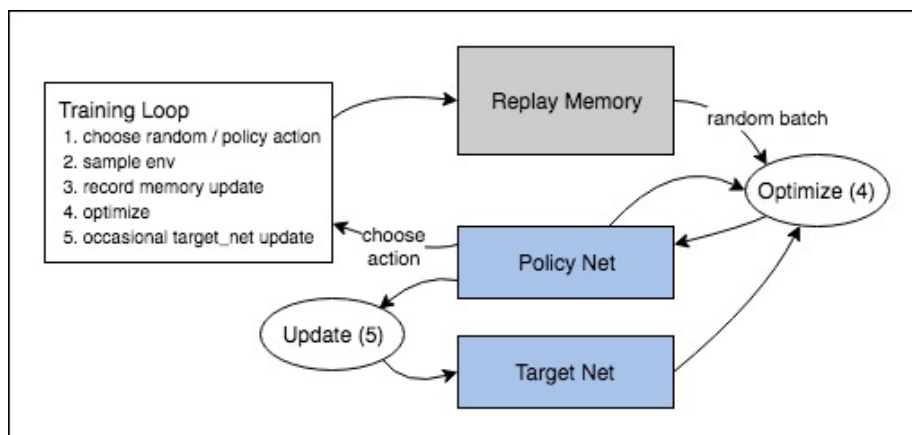


Figure 3: Training loop overview

2 OpenAI Gym Wrapper

Wrappers are used to transform an environment in a modular way. [Link to gym wrappers](#).

```
env = gym.make('PongNoFrameskip-v4')  
env = MyWrapper(env)
```

WarpFrame Wrapper

Warp frames to 84x84 as done in the Nature paper and later work. Expects inputs to be of shape height x width x number_of_channels

PyTorchFrame Wrapper

Pytorch expects images in the form [number_of_channels, height, width] hence this wrapper transforms image from [height, width, number_of_channels] to [number_of_channels, height, width]

Gym Monitor

The gym monitor wrapper will assist you in capture a recording of your agent playing the game.

```
env = gym.wrappers.Monitor(env, "path/to/store/recordings")
```